

DENNIS WOLF

INVESTIGATING EVENT SUBSCRIPTION
MECHANISMS IN BPMN

INVESTIGATING EVENT SUBSCRIPTION MECHANISMS IN BPMN

DENNIS WOLF



< Any Subtitle? >

August 2017 – version 1

Dennis Wolf: *Investigating Event Subscription
Mechanisms in BPMN*, < Any Subtitle? >, © August 2017

ABSTRACT

Business Processes have become an essential tool in organizing, documenting and executing company workflows while Event Processing can be used as a powerful tool to increase their flexibility especially in distributed scenarios. The publish-subscribe paradigm is commonly used when communicating with complex event processing platforms, nevertheless prominent process modelling notations do not specify how to handle event subscription.

At the example of BPMN 2.0, the first part of this work illustrates the need for a flexible usage of event subscription in process models and derives new requirements for process modelling notations. An assessment of the coverage of these requirements in BPMN 2.0 is presented and shortcomings are pointed out.

Based on the identified requirements, this work presents a new concept for handling event subscription in business process management solutions, predominantly built on the notion of event buffers. The concept includes an extension to the BPMN meta model, specifies the semantics and API of a new event buffering module and describes the changes necessary to the behaviour of the process engine.

For evaluation purposes, the concept has been implemented as a reusable Camunda Process Engine Plugin that interacts with the academic Complex Event Processing Platform UNICORN.

ZUSAMMENFASSUNG

Kurze Zusammenfassung des Inhaltes in deutscher Sprache...

CONTENTS

1	INTRODUCTION	1
1.1	Motivation	2
1.2	Contribution	2
1.3	Structure	2
2	BACKGROUND ON EVENT-DRIVEN BUSINESS PROCESS MANAGEMENT	3
2.1	Business Process Management and Technology	3
2.1.1	Business Process Model and Notation	3
2.1.2	The Process Engine	3
2.2	Complex Event Processing	3
3	PROBLEM STATEMENT	5
3.1	Motivating Examples	5
3.2	Event Occurrence Scenarios	9
3.3	Requirements Definition	11
4	ASSESSMENT OF CURRENT BUSINESS PROCESS MANAGEMENT SOLUTIONS	13
4.1	BPMN Models in presence of the Event Occurrence Scenarios	13
4.2	Implementation of Early Event Subscription using standard Camunda	17
4.3	Discussion	21
5	FLEXIBLE EVENT SUBSCRIPTION IN BPMN	23
5.1	BPMN Extension	23
5.1.1	Adding basic subscription information	24
5.1.2	The time of event subscription modeled in BPMN	25
5.1.3	Using Process Variables in Event Queries	28
5.1.4	Advanced Buffer Parameters	28
5.2	Design Decisions	30
6	AUTOMATIC SUBSCRIPTION HANDLING	33
6.1	Buffered Event Processing	33
6.2	Extended Process Engine Behavior	36
7	REFERENCE IMPLEMENTATION	39
7.1	Extending the Event Processing Platform Unicorn	39
7.2	Event Subscription Handling in Camunda	39
8	RELATED WORK	41
9	CONCLUSIONS	43
	BIBLIOGRAPHY	47

LIST OF FIGURES

Figure 1	BPMN Model of a Logistics Process using events for route optimization (Example 1.1)	5	
Figure 2	Transport via English Channel that is timed to a delivery slot (Example 1.2)	6	
Figure 3	Model of a retail order management process (Example 2)	7	
Figure 4	Possible event occurrence times in relation to a process execution life cycle	9	
Figure 5	Abstract Process using an Intermediate Catch Event	13	
Figure 6	Standard Intermediate Catch Event	14	
Figure 7	Intermediate Event with a parallel Timer Event	14	
Figure 8	Event Element in parallel process flow	15	
Figure 9	Event Buffering through an auxiliary Buffering Process	16	
Figure 10	Generic Example Process in Camunda for Occurrence Scenario <i>O3</i>	18	
Figure 11	Generic Example Process in Camunda for Occurrence Scenarios <i>O4</i> and <i>O5</i>	18	
Figure 12	Auxiliary Buffering Process in the Camunda Modeler	19	
Figure 13	Auxiliary Event Delivery Process in Camunda Modeler	20	
Figure 14	The concept for flexible event subscription involves three modules: A BPMN extension, enhanced process engine behavior and buffered event handling.	23	
Figure 15	Flight Booking process using a consuming buffer	30	
Figure 16	Shared consuming buffer in Complaints Handling	30	

LIST OF TABLES

LISTINGS

ACRONYMS

INTRODUCTION

Given the increasing competition on the global market place, companies are seeking to improve their value and reduce costs. In many areas, Business Process Management has been chosen as one of the tools to help stay competitive. Especially large enterprises, but also small to medium businesses formalize their workflows in business process models to allow archiving, documentation and automatized management and execution.

With Business Process Technology in constant progression, the opportunities that the field has to offer are ever growing. Since the recent years, many efforts have been dedicated towards bringing together business processes and Complex Event Processing. By the help of Event Processing Systems, companies are trying to get a hold of the exponentially growing amounts of data that occur in today's IT environment. Incorporated in process executions, events heavily increase their capabilities and flexibility. They can be utilized for intraorganizational communication between processes or business departments, but also allow to respond to external situations within seconds or milliseconds. The Business Process Model and Notation (BPMN), an industry standard for representing business processes both visually and textually, natively supports the use of events in a plethora of ways. Events are considered a main building block of a feature-rich process modeling language and can, for instance, be used for instantiating processes, communicating between process participants or to support decisions.

An interaction with a Complex Event Processing platform generally follows the publish-subscribe paradigm: The event consumer contacts the platform and issues a subscription to a specific subset of available events. An event producer, for example a Vehicle providing its current GPS location, publishes information to the event processing platform, which is then forwarded to every consumer that had subscribed. Intermediate Events are a basic way to implement event-based communication in BPMN and facilitate the reception of external message events. Nevertheless, the BPMN specification does not specifically consider the publish-subscribe workflow and offers only limited flexibility when it comes to incorporating event subscription and un-subscription operations in business process models. The consequences and possible solutions to this lack of specification are investigated in this work.

1.1 MOTIVATION

- distributed -> hard to control - subscription before occurrence time
- subscr. not considered in detail in bpmn - other research towards
..., - increasing demand for using events in processes - problems can
occur when mis-used, delay or blockage, significant time and finan-
cial loss -> brief, textual example - a more flexible approach to event
subscription might help resolve those issues

1.2 CONTRIBUTION

(o) reviewing the problem from different perspectives, deriving re-
quirements (1) assessment of the capabilities of standard bpmn (2)
Proposition of a BPMN extension for flexible event subscription, its
advantages (3) derived requirements to process engines and CEP
platforms (4) A reference implementation using Camunda and UNI-
CORN

1.3 STRUCTURE

write

BACKGROUND ON EVENT-DRIVEN BUSINESS PROCESS MANAGEMENT

2.1 BUSINESS PROCESS MANAGEMENT AND TECHNOLOGY

- Brief: what is a business process -> definition - activity / instance life cycle
 - What's the management process, brief life cycle

2.1.1 *Business Process Model and Notation*

- modeling notation, focus on ..., 2.0 - main elements: -> briefly explain
- events: why? how? - Message - choreographies

2.1.2 *The Process Engine*

- What's the infrastructure to manage and execute business processes
- what does the engine do. - Introduce Camunda

2.2 COMPLEX EVENT PROCESSING

- what does it do - how does it generally work -> stream processing and | input > processing > output | - Event queries: Esper - event producer, consumer, subscription -> pub/sub - interplay BPT to CEP platforms, sub in BPMN - putting cep queries into bpmn models => heiko's thesis or other related work?

PROBLEM STATEMENT

This section will further define the problem and derive formal requirements to event subscription mechanisms

illustrate the complexity through 2 sample scenarios that are similar but different

3.1 MOTIVATING EXAMPLES

To allow a better understanding of the issue, event-driven use-cases from two different domains are presented in the following. The cases are revealed through their standard BPMN representation, which can cause problems in certain real-life situations. It is illustrated, why the time of event subscription is of great importance and motivate to study the mechanics and implications of event subscription in business processes.

DELAY OF A LOGISTICS PROCESS The first example([Figure 1](#)) is taken from the logistics domain and shows a truck transport that has to cross the English Channel. The truck driver receives the transport plan for his next tour from France to the UK. By default, the company crosses the Channel using the Eurotunnel, an underground train connection between London and Paris.

After loading the goods at the factory, the truck will head towards the check-in location of the Eurotunnel. If everything runs on schedule, the truck will cross the channel on the train and then deliver the goods in Great Britain. Alternatively, the process considers a route using the ferry from Calais(FR) to Dover(UK). The Eurotunnel administration publishes delay information approximately every 30 minutes through an RSS feed on their website. While it mostly operates on schedule, delays ranging from 15 minutes to several hours occur regularly. It can happen that new information is not published for multi-




Figure 1: BPMN Model of a Logistics Process using events for route optimization (Example 1.1)




Figure 2: Transport via English Channel that is timed to a delivery slot (Example 1.2)

ple hours. Significant delay events (delay > 30 minutes) are received through a boundary catching message event attached to the activity *Drive to Eurotunnel Check-In*. The boundary event is interrupting, hence the activity is canceled if a delay occurs. The transport continues towards the ferry terminal and crosses the English Channel over sea. After crossing the channel, the goods are delivered to the recipient.

show a cep query for that scenario to make it more precise

According to the BPMN specification, the listening for the boundary event starts as soon as the related activity is activated.

is that exactly correct?

Given that events arrive every 30 minutes, there can be a gap of up to half an hour, before the first information becomes available. In the worst cases, when data isn't published for several hours, this gap will be even bigger. Let's consider a very busy weekday. A technical fault occurred in the tunnel earlier and the train runs 3 hours behind schedule. The last information on the RSS feed was published at 2:35pm. At that time the truck driver is still in the process of loading goods, finishing the activity at 2:40pm. Following the process definition, the driver now departs towards the Eurotunnel check-in. The system publishes updated information at 3:15pm: operations are still 2:30h behind schedule. The message gets received through the process and the truck driver takes the alternative route to the ferry, but only after heading to the Eurotunnel for 35 minutes. The late change of plans causes an unnecessary delay to the shipment.

Figure 2 is an extension of the the transport process. In logistics, it is common that a delivery cannot be accepted at arbitrary time. Instead, the receiving party assigns delivery windows to the transport company. The transport must arrive during the given time window, otherwise the delivery cannot be completed. After crossing the English Channel, the process model shows the catching of a message event containing the desired final arrival time at the factory. There is an agreement with the factory, that the delivery slots will be approved 2 hours before the expected arrival. If the current ETA of the transport is greater or equal to the arrival time, the driver will head to




Figure 3: Model of a retail order management process (Example 2)

the drop-off point straight away. If the transport is ahead of schedule, the driver will have to delay the delivery to match the time window.

The presented process model illustrates another complexity of using events in processes. Again, the listening to the announcement of a delivery window will start when the event element is reached, in this case only after crossing the English Channel. Until an event has been received, the process will not continue. Much worse: if the receiving party sends out the arrival time information too early, while the truck is crossing the channel, the event is missed. If it is not issued again, the process cannot receive a message and will get stuck indefinitely waiting to catch the event.

Neither of the two presented catch events allow for an efficient and reliable execution of the process. They can cause unnecessary delays and even blocking of the process execution. The remainder of this work will further analyze the capabilities of BPMN to express event-usage scenarios and propose solutions to the mentioned problems.

Consider more possible event occurrence times to prepare for the next chapter

UP-TO-DATE SHIPPING INFORMATION FOR AN ORDER A similar situation can be observed in the Order Management process presented in Figure 3. It describes the interaction between customer and seller in a traditional distance retail scenario: After browsing the product catalog, the customer requests a quote for the articles he or she is willing to buy. The retailer makes an offer including an approximation of the expected shipment date and sends it to the customer. That quote is then either accepted or not and the payment is issued if necessary. Once the retailer is informed about the placement of the order, the products are packed and shipped as soon as possible. For articles that are not currently in stock, the retailer must await the shipment from the factory. If any of the factory-shipments is delayed, the re-

tailer cannot ship in time and will announce a delayed shipment date to the customer. This situation is modeled through a non-interrupting boundary event attached to the *Prepare Shipping* activity, which triggers the sending of the updated shipment date to the customer.

The process shows a number of similarities, but also differences in terms of event-use when compared to the previous process example. At first we want to look at the three intermediate catch events, *Quote received*, *Order canceled* and *Order placed*. In each of the cases, the event to be caught is the direct response to a message that was sent right before. While the process will also enter a waiting state until the response arrives, that waiting is not to be interpreted as an unnecessary delay to the process execution. Other than in Example 1.2 (Figure 2), there is nothing useful to do before the response is received. It is furthermore worth noting that the response messages cannot be missed, because the message catch event immediately follows the message send event.

A different situation holds for the boundary event *Shipment Delay*. While the subscription to a Eurotunnel event can be issued at any time, it does not depend on any process data, the shipment delay has to be observed for each product that is part of the order. A subscription can therefore not be executed before the activity *Prepare Quote* has finished executing

show a cep query for that scenario to make it more precise

By process definition, the system will listen to shipment delays once the activity *Prepare Shipping* has started executing, much later than possible. Any events that occur between the two activities, that means, for example, when the customer makes the decision about accepting or canceling the order, cannot be considered in the process execution and the customer will not be informed about a possible late shipment.

write about service interaction patterns somewhere

The two presented examples have illustrated the complexity of using events in business process, especially when the possible event occurrence times are considered in detail. Differences have been pointed out as to how exactly the event is placed in the process, if it waits for a direct response to an earlier request or if the event occurrence unrelated to the execution of that very process. Motivated by this complexity, it is the goal of this work to identify the capabilities of BPMN to handle event subscription in business processes.

this chapter must be closer related to event subscription. Don't make it too general. subscription is the thing

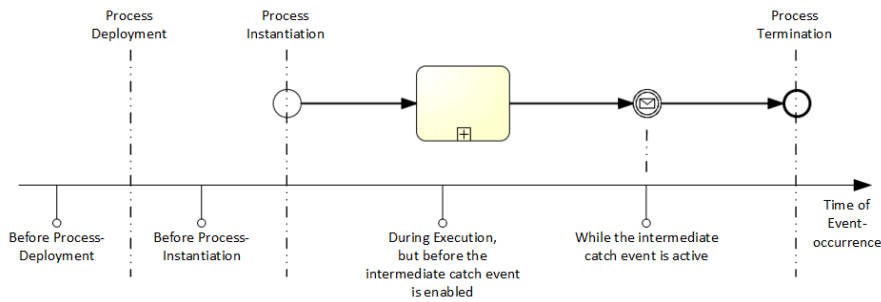


Figure 4: Possible event occurrence times in relation to a process execution life cycle

3.2 EVENT OCCURRENCE SCENARIOS

Given the motivating examples, I am deriving a generic set of event occurrence scenarios. Each of these scenarios can occur in the real world and process implementations need to be capable of handling them to avoid negative effects.

EVENT SUBSCRIPTION TIME The most important variable to consider is the time of event occurrence.

reformulate

According to the BPMN specification, it is possible to catch an event if it occurs after the event element is enabled. As shown before, it is often impossible to control occurrence time and events do occur outside of these time windows. We specify the possible event occurrence times in relation to the life cycle of a process that utilizes a BPMN Intermediate Event

ref process lifecycle

Figure 4 shows the life cycle steps of a process and an instance from the deployment of the process until the undeployment and uses a timeline to illustrate that an event might occur at any time during this cycle. More precisely, an event is always considered to occur before or after a life cycle step or in between two consecutive steps.

How about system-deployment/process engine start and process undeployment? show in illustration, but say in text that we simplify this for now. after undeployment is essentially before deployment of a new process; Before Engine start is also before pr. deployment and we presume that an engine is running and does not stop.

Given the relevant life cycle steps, *process deployment*, *process instantiation* and *Event enablement*, the following occurrence scenarios are distinguished in this work:

- A. O1 After the enabling of the BPMN event (BPMN default)

- B. O2 The event does not occur
- C. O3 Between Process instantiation and the enabling of the BPMN event
- D. O4 Between Process deployment and process instantiation
- E. O5 Before Process deployment

add a back reference to the examples. In example XY, events can occur before... whereas in example...

For a flexible and efficient use of events in business processes, it must be possible to use events that occur in any of these phases. To make sure that an event can be caught, no matter at which time during the phase it occurs, the subscription to the CEP platform must happen at the beginning of the occurrence phase. It follows that the event subscription must be possible at system start, at process deployment, at process instantiation, at any time during process execution and when the BPMN Event element is enabled.

Explain that, to catch events at these times, the subscription has to happen before. Therefore we derive event subscription times

EVENT SUBSCRIPTION DEPENDENCIES It is important to note that the subscription to an event source can depend on additional context information or process data. This can be a significant limitation to the possible subscription time.

ref to process model

shows a logistics process that uses event data about the GPS position of a certain truck to keep the estimated time of arrival of the transport updated. Whenever it receives an updated GPS position, the ETA is re-calculated; once the *arrival*-event has been received, the process finishes.

this example is not good, because we are not interested in a gps event that occurs earlier. Find an example where you would like earlier events, but subscription is not possible

Before the subscription to that specific truck gps event can happen, the process must determine the *truckId* to use in the event query. Only when the *truckId* is available, the subscription can be executed. This example illustrates how a query filter expression can depend on context data, but it might as well be the event source itself that differs depending on the particular execution.

there could be an xor gateway and following two different events and only one of them can get executed

solution would be to listen to all gps, but potentially too much data. Decision must be made cautiously! <= Where should I mention this? maybe later in the concept

3.3 REQUIREMENTS DEFINITION

The previous sections have exemplified how the execution semantic offered by the BPMN specification limits users in the use of events in business processes. Now these shortcomings are formalized into an additional set of requirements that must be met by a process execution environment to enable event handling in the extended set of event occurrence scenarios. The formal requirements will later be used to evaluate the capabilities of current Process Management Solutions ([Chapter 4](#)) and to develop a new concept to handling event subscription in business processes ([Chapter 5](#)).

R1: FLEXIBLE EVENT SUBSCRIPTION TIME

R1.1 EXPLICITNESS: For each event that is used in a business process, it must be possible to derive the time of event subscription from the process model. The time of subscription may either be explicitly stated or defined implicitly.

R1.2 FLEXIBILITY: The time of subscription can be influenced to catch events according to any of the event occurrence scenarios O1, O2, O3, O4. In other words, the process model defines the earliest acceptable time for an event occurrence to be considered in the process execution. The necessary options are *since system start*, *since process deployment*, *since process instantiation*, *from an arbitrary but explicit time during process execution*, or *since enabling of the Event Process Element*.

limited by subscription dependencies

change the options back to the times of subscription. Mention that the subscription is necessary before the time of event occurrence, but too early subscription is also a problem.

target only intermediate events. (limitation, although it also makes the most sense.)

R2: AUTOMATIC SUBSCRIPTION HANDLING

R2.1 SUBSCRIPTION: The subscription to event sources is handled implicitly by the process execution environment as defined by the process model.

R2.2 REMOVAL OF SUBSCRIPTION: The removal of a subscription from the system is handled automatically as soon as a subscription becomes unnecessary.

R3 EVENT BUFFERING: To make all events since the subscription time available during process execution, matching events need to be stored temporarily.

buffer policies?

ASSESSMENT OF CURRENT BUSINESS PROCESS MANAGEMENT SOLUTIONS

The lack of flexibility in handling event subscription in business processes has been outlined in the previous chapters and a set of extended requirements to process management solutions have been presented. In this section I take a closer look at the capabilities of current solutions with regards to the event occurrence scenarios to get a better understanding of the issues that arise when working with event subscription in business processes. The assessment will be carried out using BPMN and Camunda, a state-of-the art and widely adopted business process engine. The main goal is to identify and illustrate the shortcomings of the current process technology stack. These shortcomings will be referenced in addition to the presented requirements to develop a more refined subscription handling model in the following chapter.

"subscription handling model"?

which functionality should be evaluated exactly?: all occurrence scenarios, but no buffer policies. The buffer will always store the last version of the event and also deliver that version.

4.1 BPMN MODELS IN PRESENCE OF THE EVENT OCCURRENCE SCENARIOS

Chapter X has revealed that processes can run into deadlocks if events do not occur at the right time

.

Figure 5 shows a generalized process that uses an Intermediate Catch Event just before process termination. In this section I first describe for each Event Occurrence Scenario how this simple event implementation behaves in presence of the given scenario. I then evaluate if it is feasible to create a BPMN model that is free from deadlock in these situations.




Figure 5: Abstract Process using an Intermediate Catch Event




Figure 6: Standard Intermediate Catch Event




Figure 7: Intermediate Event with a parallel Timer Event

SCENARIO 01: THE EVENT OCCURS AFTER THE ENABLING OF THE BPMN EVENT The first scenario represents the most simple case, that is also natively supported by the BPMN 2.0 specification. When the event occurs after the Event element has been enabled, the event will be received and the process can proceed normally. The use of a standard Intermediate Catch Event does suffice to cover this situation.

SCENARIO 02: THE EVENT DOES NOT OCCUR In certain situations an event might not occur at all. Given a basic event implementation like in [Figure 5](#), the process flow will get to a halt once it reaches the Intermediate Catch Event and will not be able to proceed. While, depending on the process design, this might be the desired behavior, in many situations this is not acceptable.

Let's consider a process that is supposed to wait for approval for a certain amount of time and trigger an additional request if the approval has not been issued before the deadline. [Figure 7](#) shows how this behavior can be implemented using an Event-based Gateway which puts a Timer Event in parallel to the Intermediate Catch Event. This extension will make sure that a process does not run into a deadlock state if the expected event does not occur.

I mention an example, but that example is not exactly illustrated in the process




Figure 8: Event Element in parallel process flow

according to the spec: what exactly will happen to the active catch event once the timer fires?

SCENARIO O3: OCCURRENCE BETWEEN PROCESS INSTANTIATION AND THE ENABLING OF THE BPMN EVENT In case the event occurs during process execution, but before the BPMN event element is enabled and thus listening for events, the occurrence will not be considered in the execution. The process will get stuck at the Event Process Element as if the event did not happen at all. To avoid a deadlock in this scenario, a solution is to execute the Intermediate Catch Event in parallel to the rest of the process flow using a Parallel Gateway. This is illustrated in [Figure 8](#). The time of subscription to the event can be controlled by the position of the parallel split: To implement an event subscription right after process instantiation, the Parallel Gateway has to be the first element after the Start Event (that means *Part 1* in the illustration is empty). To implement event subscription at a specific point during process execution, part of the process can execute before reaching the Parallel Gateway. In [Figure 8](#), the event may occur at any time during the execution of the collapsed sub-process *Part 2*.

SCENARIOS O4 AND O5: BEFORE PROCESS INSTANTIATION Any Events that happen before process instantiation will not be considered in a standard Intermediate Catch Event. That applies to both scenarios, the occurrence between deployment and instantiation (O4) and an occurrence time before the deployment of the process in the Process Engine (O5).

To create a Process Model that allows to catch an event before the process instance exists, three new elements are introduced: (1) An additional *Auxiliary Buffering Process* that can catch an incoming event, (2) an *Event Buffer*, a temporary data-store that keeps event data until required by the *Original Process*, (3) an *Auxiliary Event Delivery Process*, that retrieves events from the buffer and makes them available to the *Original Process*. [Figure 9](#) reveals the interaction of the *Original Process*, the two auxiliary processes and the data-store. To start listening for an event, the *Auxiliary Buffering Process* has to be instantiated through




Figure 9: Event Buffering through an auxiliary Buffering Process

the a message start event containing the information necessary for the event subscription. The process starts listening for the event and writes the received event to the temporary data-store. The given process design is able to handle multiple event occurrences, because the receiving activity is looping. The buffering process terminates once the *Unsubscribe* event is received.

how should the ABP be started? manually?

The *Original Process* can be started any time after the buffering process. In Figure 9, the Intermediate Catch Event has been explicitly split into three events: An initial Send Event to request events, a Catch Event to receive and a final Send Event to signal that no events shall be received anymore. The initial Send Event instantiates the *Auxiliary Event Delivery Process*, which tries to read from the Event Buffer and deliver the event to the *Original Process*. Once there is data available in the buffer, it is sent by the sending activity. The central looping activity will retry reading from the buffer until data becomes available and will only be terminated once the *Stop*-event occurs. The *Original Process* can receive the event using a standard Intermediate Catch Event even when the event occurs before the instantiation of *Original Process*, so it handles scenario *O4*. Moreover the *Auxiliary Buffering Process* is not bound to a specific event, it works generically with any event information that is passed on to it. For that reason it is also

not bound to a specific process deployment and can buffer events even before a process has been deployed, so it handles scenario *O5*. Given that the buffering process can alternatively be started using an explicit Message Send Event during process execution and the process does not stop listening until the Original Process has received the event, scenarios *O1* and *O3* are also supported.

what exactly is passed around | how many instances of each process | overwrite or append to buffer | this is only one solution to do this, one that requires minimal changes in the original process

4.2 IMPLEMENTATION OF EARLY EVENT SUBSCRIPTION USING STANDARD CAMUNDA

The previous chapter has shown that it is possible to create BPMN models to match each of the Event Occurrence Scenarios, though for the scenarios *O4* and *O5* the solution becomes increasingly complex. In the next step I investigate the capabilities of Camunda, a modern and actively developed Business Process Engine that is available under an open source license. Camunda shall be used without any code customization, that means as offered on the website. The solution presented for the last two scenarios has proven capable enough to handle all Event Occurrence Scenarios, therefore the goal is to implement this solution. It will be necessary to create the two auxiliary processes and a data-store in addition to the original process that makes use of the event buffering.

Two generic sample processes have been modeled for demonstration purposes. [Figure 10](#) shows a simple process with an explicit subscription activity to represent the listening to the event after process instantiation but before reaching the Catch Event (Scenario *O3*). It follows a sample activity that takes 15 seconds (implemented using a *Script Task*), the Intermediate Catch Event and another Script Task that displays the content of the received message. The example for scenarios *O4* and *O5* (see [Figure 11](#)) comprises the following elements: After the start event follows an Intermediate Catch Event, then an activity that prints the message of the event to console and last the Process End Event. Both figures show the processes as modeled in the Camunda Modeler.

AUXILIARY BUFFERING PROCESS The task of this process is to subscribe to a CEP Platform using a provided event query and start listening for events. Any incoming event must be stored in a data-store (*Event Buffer*). UNICORN, an Esper-based academic event processing platform, will be used in this example. A local MySQL database has been chosen for buffering the event data because it's freely available, quick to set up, offers standardized access via SQL queries and Java




Figure 10: Generic Example Process in Camunda for Occurrence Scenario O_3

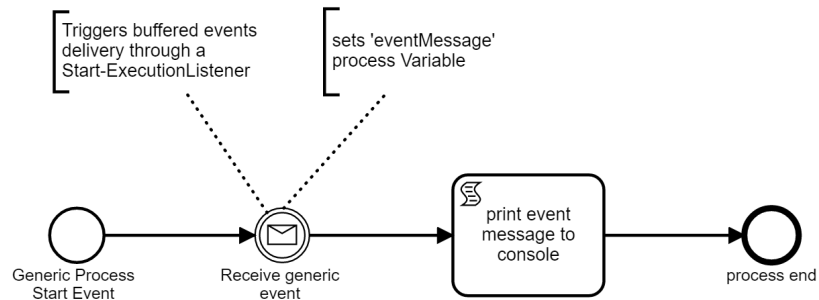


Figure 11: Generic Example Process in Camunda for Occurrence Scenarios O_4 and O_5

connectors and will persist data to the local harddrive by default. As UNICORN also requires an SQL database, the MySQL instance can be used in both cases.

Figure 12 shows the final Buffering Process modeled in the Camunda Process Modeler. The process can be instantiated by issuing a *Buffering Task* message. This message must contain three data fields: *processDefinitionId*, to know which process definition the buffered messages belong to; *messageName*, the name of the message event within the process; *query*, the event query in the Esper Query Language. Camunda will make the message data automatically available in the process instance as process variables, so they can be used during the execution of the Buffering Process. After instantiation, the process reaches the activity *Subscribe to Event Source*, a *Java Service Task* that executes a HTTP call to the UNICORN platform. That call registers the event query in UNICORN.

it tells unicorn its own instanceId, so that unicorn can correlate events to that exact instance

Afterwards, the process reaches the receiving activity *Wait for unsubscribe event* that will terminate the process as soon as the *Unsubscribe* event has been received. As long as this activity is active, events can be received through the attached Non-Interrupting Boundary




Figure 12: Auxiliary Buffering Process in the Camunda Modeler

Event. Incoming events have a field *eventBody*, which contains the event information and becomes available through a process variable with the same name.

does the message from UC have a field 'eventBody'?

The boundary event triggers the service task *Write eventBody to datastore*, which takes the data from the process variable and writes it to the MySQL Database Instance (*Global Event Buffer*).

AUXILIARY EVENT DELIVERY PROCESS The delivery process (see [Figure 13](#)) reads the latest data from the buffer and sends it to the process instance. It can be started with a message that contains the *processInstanceId* and the *processDefinitionId* of the requesting process and the *messageName* of the Message Event that is requested from the buffer. A *Delay Timer Event* has been inserted to make sure that the receiving process is already in listening state, the execution happens asynchronously. It follows the service task *Retrieve event from buffer*, which executes Java code to read from the MySQL Database *Event Buffer* and store the event information in a process variable named *eventMessage*. The content of that process variable is sent to the *Original Process* in the Send Event, afterwards the execution is finished.

INTERACTION OF THE PROCESSES In this implementation of flexible event subscription, the action of subscribing to the event source and the reception of events in the Original Process are splitted into two separate parts, each supported by an auxiliary process. To initiate the subscription at the event source, the Auxiliary Event Buffering process has to be started. For scenario *O3*, this happens through an extra activity (*Trigger Buffering Process*) during process execution, so that events after process instantiation are received by the Buffering Process. In scenarios *O4* and *O5*, the subscription and thus the instantiation of the Buffering Process must happen before the instantiation of the Target Process. As there is no such mechanism in the standard

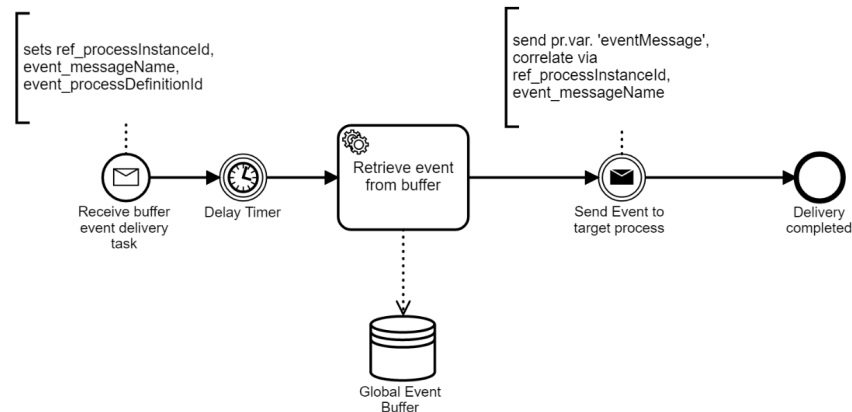


Figure 13: Auxiliary Event Delivery Process in Camunda Modeler

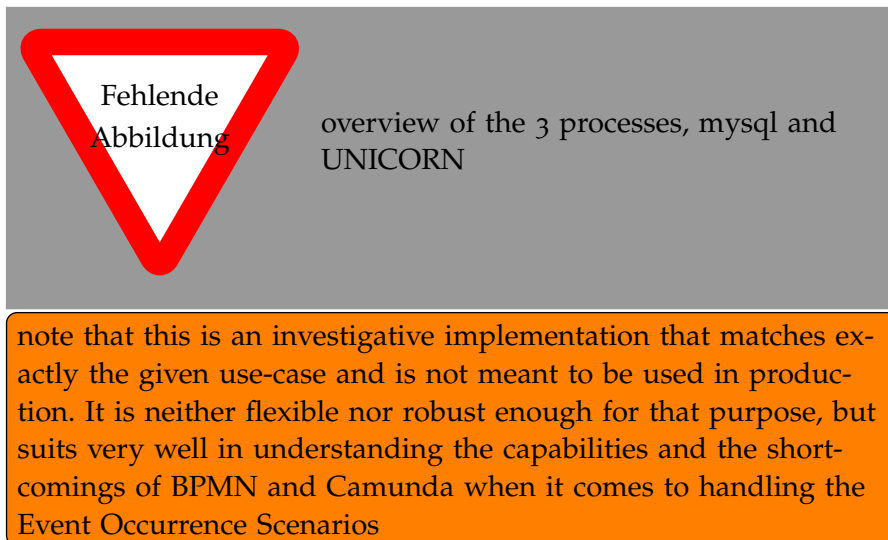
Camunda Process Engine, the Buffering Process must be started by hand, providing the *processDefinitionId*, the *messageName* and the *eventQuery*.

Now that the *Buffering Process* is running, any events matching the query will be stored to the buffer. When the Target Process reaches the Catch Event, a request for buffered events is sent as a message to trigger the *Auxiliary Event Delivery Process*. This message is sent using a short piece of Java code that gets executed when the Catch Event is reached. The code is invoked by a Start ExecutionListener attached to the Catch Event. ExecutionListeners are offered by Camunda to execute own Java programs before or after relevant events during process execution, like the execution of an element in the process. While the Original Process will now start listening for the desired events, the Event Delivery Process will send the buffered events as messages to the Original Process.

If no events have been received yet, all the involved processes remain active: the Buffering Process will keep listening for an external event. The Delivery Process will send an event to the Original Process as soon as there is one in the buffer. The Catch Event in the Original Process will keep listening for an Event.

the termination of the processes is not yet implemented in Camunda

get wording straight: Original process, target process, requesting process, main process | also always italic or never



4.3 DISCUSSION

The goal of this chapter was to get a better understanding of the capabilities of the tools when it comes to covering all event occurrence scenarios. Even though it has proven possible to implement a flexible event subscription time using standard BPMN 2.0 and Camunda, the success comes at a cost. The downsides of the presented approach are presented in the following.

It was necessary to create two generic auxiliary processes for event buffering, to connect to a MySQL data-store and use ExecutionListeners to execute custom Java code in Camunda to cover all scenarios, *O1* to *O5*.

give them short names for reference? Or make one of them Requirement 5

MISSING AUTOMATIC SUBSCRIPTION HANDLING

In the presented process models, separate process elements had to be added to handle event subscription and initiate event delivery. That conflicts with requirement *R2*, which states that the subscription and un-subscription must be automatically handled by the process engine. For the scenarios *O4* and *O5* the Buffering Process has to be triggered manually, because it must be executed before the target process is running. Camunda does not handle external event subscription itself, especially not before the process is running.

ADDITIONAL MODEL COMPLEXITY

As additional process elements have to be added to handle event subscription and delivery, the models become more complicated and are less concentrated on the business case.

there could be a requirement that states that there should be no additional process elements unless there is an explicit subscription time

BUFFERING IS AN IT TASK

The auxiliary processes are not business tasks and are thus not suited to be modeled in BPMN. Desired functionality can be put into Camunda BPMN models thanks to its flexibility to use Java code in Service Tasks or Event Listeners, but naturally the full functionality of the Event Buffer cannot be expressed using BPMN.

what exactly is the issue here?

ADDED LOAD ON THE PROCESS ENGINE

Because of the aux processes, two additional processes have to be deployed in the process engine and are potentially running in parallel to any given process instance. For each Event Element used in a process the engine has to run an instance of the Buffering Process and, eventually, an instance of the Buffer Delivery Process. That puts additional load on the process engine, which might prevent business critical processes from executing delay-free.

Even when the number of deployed and running auxiliary processes can be reduced through further optimizations there remains an event-management overhead as every event has to be handled twice: once when it is stored in the buffer and once when it's delivered to the target process.

HIDDEN PERFORMANCE LIMITATIONS OF THE PROCESS ENGINE

Given the large amount and high frequency in that events can occur in reality, optimal performance is required for an event-buffering module. Running essential parts of the buffering within the process engine might pose performance limitations that cannot be influenced without tempering with the process engine code.

Introduce a new concept. three pillars: bpmn-x, process engine behavior, event handling api

5.1 BPMN EXTENSION

this is the explanation for the process designer? or is there an additional event engineer?

Given the additional requirements and the shortcomings identified in the previous sections, the following two chapters present an extension to the BPMN event handling model. At first, an extension to the Business Process Model and Notation (BPMN) is described, which aims at providing the Process Designer with more flexible Event Handling capabilities according to Requirement *R1*. Afterwards, [Chapter 6](#) clarifies the changes necessary to the event handling platform and the process engine to cover Requirements *R2* and *R3*. While the presented concepts are kept as general as possible, they are grounded in an analysis of the Esper-based CEP Platform Unicorn and the open-source process engine Camunda.

The extensibility mechanism allows to extend standard BPMN elements with additional attributes, while maintaining a valid BPMN core. Extensions are specified through an external definitions file and can be included into a BPMN process model by reference. To allow the flexible use of event subscription in BPMN models, a number of additional attributes must be added to BPMN process models.




Figure 14: The concept for flexible event subscription involves three modules: A BPMN extension, enhanced process engine behavior and buffered event handling.

The presented extension is designed for use with intermediate catch events, boundary events and receive tasks. Start events and End events are not considered because flexible subscription time is not relevant in these cases. Each of the three elements references a BPMN *Message*, the common denominator for communication within and across business processes. Semantically, the *Message* type is most suited to be extended with subscription information.

mention that others have extended the events instead, but this proposal goes another way



A UML diagram of Message, Message-Flow, MessageEventDefinition, Catch Event, Boundary Event, ReceiveTask

explain that the Message is a common element in BPMN models. It is the main generic type used for communication, collaboration. The proposed extension to the Message type is only to be used by Message Receive situations in Intermediate Catch Event, the Boundary Catch Event and the Receive Task.

By specification, *tMessage* comprises an attribute *name*, the name of the message, and *itemRef*, the reference to a BPMN *ItemDefinition*. Additionally, it inherits all attributes from the BPMN *RootElement*. In the following, the required additional attributes will be explained one after the other. A complete list is available in [Table 1](#). The goal is to retain a stand-alone model that contains all information necessary to execute the subscription to the event source.

5.1.1 Adding basic subscription information

For a basic event subscription, an event query, the platform address and optionally authorization information of the CEP Platform is required.

ref background

It is assumed that only one CEP platform is in use, whose access information is configured centrally for the current process execution environment. Consequently, there is no need to specify these two parameters on message level. The event query instead needs to be specified for every message and is added to the model as an extension attribute *eventQuery* of type *String*, which should contain the full query as interpretable by the CEP platform.

ATTRIBUTE NAME	VALUE OPTIONS (<u>DEFAULT</u>)	OPTIONAL
eventQuery	any string	n
subscriptionTime	process-deployment, process-instantiation, <u>event-reached</u>	y
bufferPolicies	Complex Type (see below)	y
BUFFERPOLICIES		
LifetimePolicy	string in ISO time-span format OR ' <u>infinite</u> '	y
ConsumptionPolicy	<u>Reuse</u> , Bounded-Reuse(n), Consume	y
SizePolicy	int (< 1 for infinite), <u>1</u>	y
OrderPolicy	<u>FIFO</u> , LIFO	y

Table 1: Available attributes in the BPMN extension for flexible event subscription

A similar approach has been taken by X and Y, who aim at enriching BPMN models with subscription information without considering the time of subscription specifically.

Find this source; explain what they do (different)

Given this first fundamental part of the BPMN extension, it is possible to execute the subscription, but the time of subscription cannot be influenced.

5.1.2 The time of event subscription modeled in BPMN

This section specifically addresses the requirement *R1.1*, aiming to provide a flexible event subscription time to be selected for each BPMN message when designing an event-driven process. Two different tools are to be offered to support all subscription times demanded by *R1.1*: Firstly, the subscription can happen in the background. Alternatively, the subscription can be modeled explicitly as a flow-element in the process. It is the task of the process designer to elaborate the correct time of subscription necessary for her use case.

The subscription will be executed automatically by the system based on the information given in the BPMN model. Further information on the exact execution flow is provided in chapter XY.

INTRODUCING EVENT BUFFERS Any event message that occurs before reaching the event element but after the time of subscription will be kept in a buffer by the system. In its simplest version, the

buffer is of length 1, that means it stores exactly one message received from a CEP platform. It always stores the latest message. When a newer message arrives, the old one is replaced in the buffer. [Section 5.1.4](#) introduces a set of advanced buffer policies to adapt this behavior further.

By default, there is no interference between the buffers of different messages, process instances or processes. Each buffer instance will contain the latest information as if it was the only buffer in the system. Performance improvements to avoid duplicate buffer content will be managed by the system without explicit action by the user. Section ... later introduces a shared, more complex usage scenario of the event buffers.

INTERPLAY OF EVENT QUERIES AND BUFFERS Modern event query languages are feature-rich and offer a large set of expressions to filter events from incoming streams.

ref background

The introduced basic event buffer can be used in connection with any desired event query and will store the latest output of that query. These two features together suffice to implement even more complex use-cases: Query windows of length n can be used to keep multiple events in the buffer, filter expressions allow to keep a subset of all events based on their attribute values, multiple streams can be joined together. As soon as the process flow reaches an event element, the latest CEP message is retrieved from the buffer. It is not consumed, that means a second event element that references the same BPMN Message will reuse the information from the buffer. If no information is available in the buffer, the flow element will remain in the waiting state until a message is received. Then, the process flow proceeds as usual.

reorganize this chapter

SUBSCRIPTION TIME AS PART OF THE BPMN MESSAGE ELEMENT

To provide the Process Designer with a simple but powerful tool to influence the time of event subscription, a field *subscriptionTime* is added the BPMN message element. The field can take one of the following three values: *Process Deployment*, *Process Instantiation*, *Event reached*. The last option is the default option, coming closest to the BPMN specification. Note that a *subscriptionTime* set to *Event reached* will remain without effect if an explicit subscription task for the same event was executed before the event is reached.

For each of the options: Define exactly (according to BPMN spec or standard literature), when in the flow the subscription is executed.

In motivating Example Ex, it is necessary to issue the subscription as early as possible, to make sure that data is available and the process execution is not delayed.

which example to reference?


Using the BPMN extension, the use case can be implemented by defining the event query and the subscription time in the BPMN model.

show what that would look like in the example. Maybe some XML?

THE EXPLICIT SUBSCRIPTION TASK

As an alternative to specifying the subscription time using the extension field *subscriptionTime* of *tMessage*, an extension to the BPMN ServiceTask is proposed.

it follows the ideas presented in ... but through referencing the message element



The extended task is used to execute the subscription explicitly as part of the process flow. A field *messageId* is added to the service task to establish a reference between the activity and the message definition. As introduced in section [Section 5.1.1](#), the extended BPMN Message definition contains the information necessary to issue the subscription to an event source. Once the Explicit Subscription Task is activated, the subscription for the referenced message is issued.

Modeling the event subscription in an explicit task can be necessary when the subscription depends on the result of another activity. In that case, the subscription cannot be issued on process instantiation, because the necessary information is not yet available. Instead, an early subscription can be implemented using the extended service task. Apart from this particular use case, the explicit subscription task enables the Process Designer to place the subscription flexibly in the process flow and give her full control over the time of subscription.

As an improvement to the options for subscription time, there could be an option "ASAP", so that the process engine issues the subscription automatically as soon as the required process data becomes available

If both tools, the extension field *subscriptionTime* and the explicit subscription task, are used for a single BPMN message, the earlier subscription of the two will be executed, the second subscription will have no effect. That means for example if the *subscriptionTime* is set to *event reached* and an explicit subscription task is inserted before the event element, then the subscription will be executed at the time the explicit subscription task is active. If *subscriptionTime* is set to *Process Deployment*, then the subscription will happen at that time and the explicit subscription task will remain without effect. In case neither of the two is used, the system falls back to the BPMN default and executes the subscription when the event element is reached.

5.1.3 Using Process Variables in Event Queries

this could be added in the requirements and referenced

As shown in example Z, it can be the case that the values of process variables shall be dynamically used in an event query. Therefore, the name of the process variable should be part of the event query. At the time of subscription, the mentioned variable is dynamically replaced by its current value. The exact notation for including process variables in event queries can vary depending on the applied query language as it may not interfere with any existing notation schemes. For the use with the Esper query language, the following is suggested: The exact name of the variable has to be surrounded by curly brackets and preceded by a # character: `#{VARIABLENAME}`. This notation is inspired by the usage of substitution parameters in SQL queries that are embedded in Esper. They take the form `$(expression)`.

reference esper docs 5.13.1. Joining SQL Query Results

In the example, the process uses the latest GPS position for a certain truck. The truck is identified by its unique ID which is part of the query: `SELECT lat, lng from GPSUPDATE where truckid = #{truckid}`.

missingref: dependent example

The use of dynamic process variable values introduces an additional complexity: Depending on the time of event subscription, the value of the process variable might not yet be available.

What does this mean for the process designer? A model that can take a state in that a subscription shall be issued, though the data is unavailable, is invalid. When will an error occur?

5.1.4 Advanced Buffer Parameters

Inspired by Mandal et al., a number of advanced buffer attributes are available through an extension attribute *bufferPolicies*.

add ref to sankalitas paper; do I elaborate on the changes made in comparison to the paper? + the reasons?

LIFE-TIME OF BUFFERED EVENTS

The *LifetimePolicy* allows to specify after which timespan elements in the buffer should be deleted. Timespans shall be defined using ISO timespan format.

ref external

The default value is *infinite*. Example A has been implemented by setting the *subscriptionTime* to *Process Deployment*, which means that there can be an infinite time difference between the action of subscribing to the event source and the reaching of the event element in one of the instances.

ref example

In case events are not published in a longer time, for example due to technical fault at the event producer, the buffer will contain older events that might not be relevant anymore. Using the *LifetimePolicy*, the process designer can express, that events should be deleted from the buffer after a certain period of time and thus avoid outdated information. The buffer is maintained automatically by the system. That of course comes at the price that the process has to remain in waiting state until a new event message arrives.

CONSUMPTION BEHAVIOR

for the reader it would not be clear, what a 'buffer instance' is

So far, the event buffers can be used isolated from each other. There is no interference between buffer instances and events are not removed from the buffer after retrieval. While for most use-cases this behavior is sufficient, more detailed control over the buffer can be desirable when a given message shall be used multiple times. Not always is it wanted, that events remain in the buffer after retrieval. An additional parameter *ConsumptionPolicy* is introduced which can take the values *Consume*, *Reuse*(default) and *Bounded Reuse*(*n*). While *Reuse* denotes the behavior that is already known, *Bounded Reuse*(*n*) will allow an element to be retrieved exactly *n* times. *n* has to be replaced by an integer value greater 0. The option *Consume* will remove an element from the buffer immediately after it has been retrieved for the first time, it is therefor equivalent to *Bounded Reuse*(1).

- given the option to consume from the buffer, it will now make a difference if the same buffer is accessed multiple times. - there are two scenarios to access the same buffer: (1) multiple times in the same instance, (2) multiple times because of parallel instances, (3) multiple times because of a shared buffer across processes - before proceeding, we need to be clear about the buffer scope: it depends on the time




Figure 15: Flight Booking process using a consuming buffer




Figure 16: Shared consuming buffer in Complaints Handling

of subscription. (1) after instantiation: buffer only instance-wide; (2) on pr depl: buffer reused across all instances; (3) system start: buffer reused across all processes

BUFFER SIZE AND ORDER POLICY

- two additional buffer policies: *SizePolicy*, *OrderPolicy*, default values

write text for given keypoints

- what if messages are defined in different ways? => I want to reuse the buffer between processes, but the message definition (especially cep query) are not the same. how will the system behave?

5.2 DESIGN DECISIONS

The target functionality of the BPMN extension was clearly defined by the identified requirements. To implement that functionality, there were a number of options to consider and design decisions to make. This chapter provides background information on the decisions that have influenced the presented concept for flexible event subscription.

Talk about an alternative solution? Table of events for a certain event source. The buffer is already there for me to pick from when designing the process.

THE TIME OF SUBSCRIPTION IS A QUESTION OF PROCESS DESIGN
-> process designer - this is why the bpmn extension is presented first
- hide complexity from the designer

THE ACTUAL BUFFER IS MOSTLY HIDDEN FROM THE USER - Buffers are implicitly defined through the BPMN model - keep the look and feel of the message catch event - minimal changes to existing models, backwards compatibility

AVOID ADDITIONAL USER INTERFACES - the user will only use the bpmn model - we don't want any other element because of complexity - the process should be self-contained, contain all necessary information for subscription and buffering. -> single point of contact

BUFFERS ARE CLOSELY LINKED TO PROCESS MODELS Messages are only buffered as soon as they are explicitly required by a model - we don't just buffer n messages because we might need them in the future. That would be a fuzzy, incalculable performance overhead. - instead we keep as little as possible in the buffer

THE BPMN EXTENSION IS BASED ON THE MESSAGE ELEMENT - if I want to talk about related work that goes another way - why are the policies a parameter of the message and not the catch event element?

REUSE EXISTING TECHNOLOGY -> we assume that a cep is present and that basic features of event queries can be used - if not present, then only very basic buffer functionality is available - but we don't want to start designing another event processing layer with duplicated functionality

AUTOMATIC SUBSCRIPTION HANDLING

After defining the functionality provided to the user of flexible event subscription, this chapter describes the changes necessary to the software infrastructure that is used for event-driven business process management. The concept requires that all subscription and event handling is executed by the system itself, without further interaction by the user. All necessary information for that purpose is provided by the BPMN model.

As described in REF, an event-driven process management setup primarily consists of ...

missingref

Changes are necessary to both, the Event Processing Module and the Process Engine. This chapter attempts to keep the change descriptions general so they can be applied to any common process engine and event processing platform. The first section describes the necessary extension to the event processing to support early subscription and event buffering. The following [Section 6.2](#) specifies the changes necessary to the behavior of the process engine as the connecting element between the BPMN model and the event processing platform.

6.1 BUFFERED EVENT PROCESSING

When reduced to the basics, a standard event processing platform works as follows: The user subscribes to events providing an event query and a notification-path. The platform responds with a unique identifier for that subscription. Whenever an event occurs that matches the provided query, the platform issues a notification to the notification-path. Subscriptions can be deleted through their unique identifier. These two operations, *subscribe* and *unsubscribe*, make the fundamental API of a CEP platform.

ref

EVALUATION OF COMMON CEP PLATFORMS - looking at the new bpmn extension reveals that a different behavior is required: notifications need to be kept in a buffer until they are requested by an entity - three event platforms were studied to check if the required functionality can be implemented natively: wso2, esper and ... - wso2: ? - esper: - ? : ? - a window will achieve something similar, but not exactly the same - moreover, it is desired that the same full-featured event queries as before can be used, no restrictions. - as functional-

ity wasn't available in any of the three, it was decided to specify an extended api.

Maybe: we consider these two operations as given, because common event processing platforms have these in common, but there is not common event buffering concept yet. even though e.g. es-per has something which goes in that direction: output clauses

AN API FOR BUFFERED EVENT PROCESSING The novel BPMN extension for flexible event subscription allows to issue a subscription for buffering well before the events ought to be delivered via the notification-path. The introduction of an event buffer as a separate entity between the varying list of notification-recipients and the event query makes an extension of the API necessary. Firstly, the *subscribe* operation has to be divided into two steps:

A. *registerQuery(queryString[, bufferPolicies]): queryId*

The call registers an event query in the CEP platform and instantiates a buffer. Matching events will be held in the buffer according to the specified policies. It returns a unique identifier to that new query registration and hence for the connected buffer. That identifier must be used to modify the query later.

bufferPolicies is an optional parameter which is provided as an object with four possible fields: *LifetimePolicy*, *ConsumptionPolicy*, *SizePolicy*, *OrderPolicy*. Refer to [Section 5.1.4](#) for a detailed specification of the semantics of the parameters. If *bufferPolicies* is not or only partly specified, the system should fall back to the default values.

B. *requestEvents(queryId, notificationPath)*

Initiates the delivery of notifications for a given *queryId* to a notification recipient. The recipient is specified through the *notificationPath*, the full address of the entity that is supposed to receive the message. Notifications are delivered asynchronously as soon as they are available. If the buffer is not empty, a message will be sent right after the *requestEvents* call. A similar operation, *addNotificationRecipient*, is available in existing CEP platforms. The difference is in the delivery of the first buffered message: *requestEvents* sends out the message from the buffer, *addNotificationRecipient* will send out notifications only for future query output.

be clearer about notificationPath. Specify in background and reference. also about addNotificationRecipient

add table with buffer policies, their possible values and the default val

improve explanations

A similar situation holds for the un-subscribe operation: Traditionally, a subscription is canceled through a unique identifier that is obtained as a result of the subscribe operation. After cancellation, no more notifications are delivered, the query is removed from the system. Given flexible event subscription, this operation must be split into two parts as well:

c. *unsubscribe(queryId, notificationPath)*

Removes a notification-recipient for a given query-id. Note that the buffer and query instance remain intact, so that other recipients can still subscribe.


d. *removeQuery(queryId)*

Completely deletes the query and its buffer, so that no notifications are sent out any longer.

All four methods are required to execute a subscription on process deployment. The query must be registered using *registerQuery* before the process instance, i. e. notification-path, is available. For each process instance, events can be requested individually using *requestEvents* and thereafter, the notification-recipient can be removed with *unsubscribe*. The query and its buffer will remain active even after any single instance has terminated. When the process gets un-deployed, the query can be deleted using *remove query*. [Section 6.2](#) describes the steps in detail.

show the steps with sample data from one of the examples

- to achieve the default behavior, each two steps have to be executed just after the other.



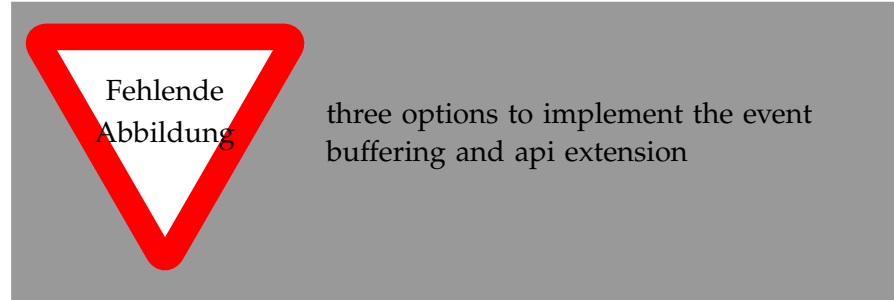
maybe provide a swagger definition for this?

what kind of API should this be? REST? Java? none specifically, but the reader might want clarification when reading this.

DIFFERENT OPTIONS TO IMPLEMENT THE API EXTENSION That extension can either be implemented by adopting the CEP platform itself, by implementing a separate middleware between process engine

and CEP platform, or by implementing a buffering module as part of the process engine. Which of the three options suits best has to be evaluated for the given use-case and existing infrastructure. In some cases it might not be possible to adapt the code of process engine or event processing platform, which leaves a separate middleware as the only choice.

Generally, extending the event processing platform is advisable?
Or put up brief pros and cons of each of the options



A reference implementation for an extended complex event processing platform is presented in [Section 7.1](#) at the example of the Esper-based CEP platform *Unicorn*. It also explains, why extending the event platform was the preferred choice in the given scenario.

PERFORMANCE CONSIDERATIONS - performance improvements through shared windows - due to its performance optimizations, an extension of the cep itself would make perfect sense - given the extended api, it is now possible to implement early event subscription from the process engine

write

6.2 EXTENDED PROCESS ENGINE BEHAVIOR

It is the task of the Business Process Engine to interpret and execute process models and connect to an event processing platform in event-driven setups. From the three relevant parts, two have already been defined, the BPMN extension and the buffered event processing module. Out of the box, a process engine like Camunda will ignore any proprietary BPMN extensions and the subscription to an event source must be especially implemented. An example for such an implementation is provided in [Section 4.2](#). One goal of this work is to automatize the handling of event subscriptions solely based on the information available through the extended BPMN model. Additional process elements should not be required. This section will clarify, which operations need to be executed by the process engine to enable the automatic subscription handling. [Section 7.2](#) demonstrates the implementation of automatic subscription handling at the example of Camunda.

PARSING ADDITIONAL INFORMATION FROM THE BPMN MODEL

It is required that the process engine is able to read the additional information from the BPMN extension (see [Section 5.1](#)) so that it is available during process deployment and execution. This affects the BPMN message element, which can contain the additional attributes *eventQuery*, *subscriptionTime* and *bufferPolicies*. Secondly, the *Explicit Subscription Task* has to be processed. It contains a reference to a Message entity within the same model. The process engine might have to be adopted to read all relevant data from the extended model.

MANAGING SUBSCRIPTION AND UN-SUBSCRIPTION As defined in the BPMN extension for flexible event subscription, the action of subscribing to an event source can happen at different times during process deployment and execution. The options and the implicit timing of subscription and un-subscription are specified in [Section 5.1.2](#). The process engine must communicate with the process engine using the four calls *registerQuery*, *requestEvents*, *unsubscribe* and *deleteQuery*, that were presented in the previous chapter. For each possible subscription time, the following briefly enumerates which operations must be executed when.

IN EVERY CASE: The return-value of *registerQuery*, a unique identifier of that query, must be stored for the related BPMN *Message*. The id is later necessary to execute the other three API methods. When an event element is reached, a call to *requestEvents* must be issued. When the execution of that event element is finished, call *unsubscribe*.

SUBSCR. ON PROCESS DEPLOYMENT: When a process gets deployed, the process engine must check if subscription information is in the model. For every *Message* element that is set as *subscribe on pr. deployment*, a call to *registerQuery* must be issued as part of the deployment process. A call to *deleteQuery* is executed when the process gets un-deployed for the same messages.

SUBSCR. ON PROCESS INSTANTIATION: When a process gets instantiated, *registerQuery* must be executed for each *Message* that is set to *subscribe on pr. instantiation*. *deleteQuery* can be called when the last reachable event element for a *Message* has finished executing or no connected event can be reached anymore. The deletion happens at the latest when the process instance terminates.

SUBSCR. THROUGH EXPLICIT SUBSCRIPTION TASK: If the control flow reaches a subscription task, the process engine executes *registerQuery* for the referenced *Message*. The execution of *deleteQuery* follows the same rules as in the preceding case.

SUBSCR. WHEN THE EVENT ELEMENT IS REACHED: Once the event element is reached, *registerQuery* must be executed for any *Message* that is not covered by one of the prior cases. *deleteQuery* must be called when the event element is finished.

be more precise about the time the calls should be executed (if possible). "reached"? "completed"? use the right bpmn words

Do I want to write about extended validity checks? soundness? The question would be: What happens if the model erroneous w.r.t. to the bpmn extension?

- handling subscription dependencies: - when are process variables replaced by their actual values - The use of dynamic process variable values introduces an additional complexity: Depending on the time of event subscription, the value of the process variable might not yet be available. - reference BPMN data elements: process INSTANCE variable -> the variable value might only be available during instance execution -> can we find an exact definition of this in the spec? - see BPMN2 spec pp.211+ : Process and Activity can have DataInput and DataOutput. DataInput can have an 'optional' attribute - during execution the variable data might or might not be available. Related Work: Francesca? -> too complex, we need a simplification for this. - what happens if the data is not available?

Time of un-subscription also must be clarified in bpmnx

did i cover all requirements?

REFERENCE IMPLEMENTATION

from [Section 5.1.1](#)

that means that in the implementation we need additional configuration values -> implementation chapter

7.1 EXTENDING THE EVENT PROCESSING PLATFORM UNICORN

7.2 EVENT SUBSCRIPTION HANDLING IN CAMUNDA

RELATED WORK

CONCLUSIONS

complete xml example of a BPMN model using the extensions

DECLARATION

Put your declaration here.

Potsdam, August 2017

Dennis Wolf

COLOPHON

This document was typeset using the typographical look-and-feel classicthesis developed by André Miede. The style was inspired by Robert Bringhurst's seminal book on typography "*The Elements of Typographic Style*". classicthesis is available for both \LaTeX and \LyX :

<https://bitbucket.org/amiede/classicthesis/>

Happy users of classicthesis usually send a real postcard to the author, a collection of postcards received so far is featured here:

<http://postcards.miede.de/>