

DENNIS WOLF

INVESTIGATING EVENT SUBSCRIPTION
MECHANISMS IN BPMN

INVESTIGATING EVENT SUBSCRIPTION MECHANISMS IN BPMN

DENNIS WOLF



Digital Engineering • Universität Potsdam

< Any Subtitle? >

August 2017 – version 1

Dennis Wolf: *Investigating Event Subscription
Mechanisms in BPMN*, < Any Subtitle? >, © August 2017

ABSTRACT

Business Processes have become an essential tool in organizing, documenting and executing company workflows while Event Processing can be used as a powerful tool to increase their flexibility especially in distributed scenarios. The publish-subscribe paradigm is commonly used when communicating with complex event processing platforms, nevertheless prominent process modelling notations do not specify how to handle event subscription.

At the example of BPMN 2.0, the first part of this work illustrates the need for a flexible usage of event subscription in process models and derives new requirements for process modelling notations. An assessment of the coverage of these requirements in BPMN 2.0 is presented and shortcomings are pointed out.

Based on the identified requirements, this work presents a new concept for handling event subscription in business process management solutions, predominantly built on the notion of event buffers. The concept includes an extension to the BPMN meta model, specifies the semantics and API of a new event buffering module and describes the changes necessary to the behaviour of the process engine.

For evaluation purposes, the concept has been implemented as a reusable Camunda Process Engine Plugin that interacts with the academic Complex Event Processing Platform UNICORN.

ZUSAMMENFASSUNG

Kurze Zusammenfassung des Inhaltes in deutscher Sprache...

CONTENTS

1	INTRODUCTION	1
2	BACKGROUND	3
3	PROBLEM STATEMENT	5
3.1	Motivating Examples	5
3.2	Event Occurrence Scenarios	5
3.3	Requirements Definition	7
4	ASSESSMENT OF THE CURRENT BP MANAGEMENT STACK	9
4.1	EVALUATING BPMN 2.0 AGAINST THE SCENARIOS	9
4.2	IMPLEMENTING EARLY EVENT SUBSCRIPTION USING STANDARD CAMUNDA	9
4.3	DISCUSSION	9
5	FLEXIBLE EVENT SUBSCRIPTION	11
5.1	BPMN Extension	11
5.2	Buffered Event Handling	11
5.3	Extended Process Engine Behaviour	12
	BIBLIOGRAPHY	13

LIST OF FIGURES

Figure 1	Possible event occurrence times in relation to a process execution life cycle	6
----------	-------------------------------------------------------------------------------	---

LIST OF TABLES

LISTINGS

ACRONYMS

INTRODUCTION

PROBLEM STATEMENT

This section will further define the problem and derive formal requirements to event subscription mechanisms

3.1 MOTIVATING EXAMPLES

- one example for independent. The subscription does not depend on a prior process result, the subscription can be done even before process instantiation
- one example for a process that uses an intermediate event that depends (subscription-wise) on the result of a previous step in the process.
 - ==> If the event occurs at a certain time, the process gets delayed unnecessarily or even run into a deadlock

3.2 EVENT OCCURRENCE SCENARIOS

Given the motivating examples, I am deriving a generic set of event occurrence scenarios. Each of these scenarios can occur in the real world and process implementations need to be capable of handling them to avoid negative effects.

TIME OF EVENT OCCURRENCE The most important variable to consider is the time of event occurrence. According to the BPMN specification, it is possible to catch an event if it occurs after the event element is enabled. As shown before, it is often impossible to control occurrence time and events do occur outside of these time windows. We specify the possible event occurrence times in relation to the life cycle of a process that utilizes a BPMN Intermediate Event

ref process lifecycle

Figure 1 shows the life cycle steps of a process and an instance from the deployment of the process until the undeployment and uses a timeline to illustrate that an event might occur at any time during this cycle. More precisely, an event is always considered to occur before or after a life cycle step or in between two consecutive steps.

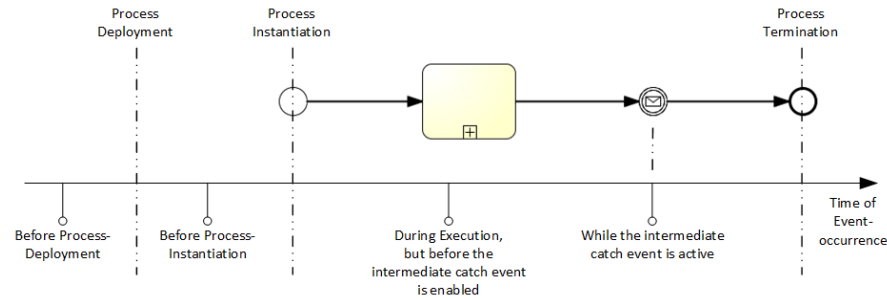


Figure 1: Possible event occurrence times in relation to a process execution life cycle

How about system-deployment/process engine start and process undeployment? show in illustration, but say in text that we simplify this for now. after undeployment is essentially before deployment of a new process; Before Engine start is also before pr. deployment and we presume that an engine is running and does not stop.

Given the relevant life cycle steps, *process deployment*, *process instantiation* and *Event enablement*, the following four occurrence phases are distinguished in this work:

- A. O1 Before Process deployment
- B. O2 Between Process deployment and process instantiation
- C. O3 Between Process instantiation and the enabling of the BPMN event
- D. O4 After the enabling of the BPMN event

add a back reference to the examples? In example XY, events can occur before... whereas in example...

For a flexible and efficient use of events in business processes, it must be possible to use events that occur in any of these phases. To make sure that an event can be caught, no matter at which time during the phase it occurs, the subscription to the CEP platform must happen at the beginning of the occurrence phase. It follows that the event subscription must be possible at system start, at process deployment, at process instantiation, at any time during process execution and when the BPMN Event element is enabled.

EVENT SUBSCRIPTION DEPENDENCIES It is important to note that the subscription to an event source can depend on additional context information or process data. This can be a severe limitation to the possible subscription time.

ref to process model

shows a logistics process that uses event data about the GPS position of a certain truck to keep the estimated time of arrival of the transport updated. Whenever it receives an updated GPS position, the ETA is re-calculated; once the *arrival*-event has been received, the process finishes.

this example is not good, because we are not interested in a gps event that occurs earlier. Find an example where you would like earlier events, but subscription is not possible

Before the subscription to that specific truck gps event can happen, the process must determine the *truckId* to use in the event query. Only when the *truckId* is available, the subscription can be executed. This example illustrates how a query filter expression can depend on context data, but it might as well be the event source itself that differs depending on the particular execution.

there could be an xor gateway and following two different events and only one of them can get executed

solution would be to listen to all gps, but potentially too much data. Decision must be made cautiously! <= Where should I mention this? maybe later in the concept

3.3 REQUIREMENTS DEFINITION

Derive formal requirements, define, make them measurable.

R1: Flexible Event Subscription Time:

R1.1: Explicitness > for each event that is used in a business process, the time of subscription must be clearly defined in relation to the process execution lifecycle > The definition is done in the process model > explicit about the event platform to subscribe to > ? What's the granularity?

R1.2: Flexibility > The subscription time can be chosen at least from the following options: [see scenarios] > Options are limited when the subscription depends on data from previous execution steps

R2: Automatic Subscription Handling

R2.1: Subscription > The subscription to event sources is handled implicitly during process deployment and execution > It is handled according to the modeled subscription time.

R2.2: Unsubscription > The unsubscription from an event source is handled automatically as soon as a subscription becomes unnecessary.

R3: Event Buffering > All events since the subscription time are available to the process. > A single buffer entity can be accessed from different process elements, process instances and processes within the environment > Buffer policies?

ASSESSMENT OF THE CURRENT BP MANAGEMENT STACK

What is the goal of this chapter and how does it fit into the structure?
Working with the Event Subscription Scenarios.

4.1 EVALUATING BPMN 2.0 AGAINST THE SCENARIOS

(BPMN 2.0 without extensions) For each Scenario: Is it feasible to create a BPMN model that is sound in presence of the given scenario? Define Soundness What is a possible example for this scenario? Provide the BPMN diagram if available. (4 pages)

4.2 IMPLEMENTING EARLY EVENT SUBSCRIPTION USING STANDARD CAMUNDA

Is it possible to implement this using out-of-the-box Camunda? Which aspects cannot be (sufficiently) implemented? How can ... be implemented in Camunda? Show details. Diagrams in appendix. (3 pages)

4.3 DISCUSSION

What are the shortcomings when using out-of-the-box business process solutions to implement Early Event Subscription? What can be implemented without problems? (2 pages)

FLEXIBLE EVENT SUBSCRIPTION

Present an abstract framework for flexible event subscription. > Including: Model <> Process Engine <> Buffer <> CEP > How does event subscription currently affect the workflow? > What should a workflow look like that allows early event subscription? > What must be explicitly stated by the user? What should be done automatically in the background? (1 page)

5.1 BPMN EXTENSION

To fulfill requirements R1.1 and R1.2, additional information has to be included in the BPMN model. By default, a BPMN intermediate event does not have information on the time of subscription or the event query. The BPMN specification offers BPMN-X extensions to add custom properties or elements to a model.

To accomodate the required information, the following extension is proposed: > The extension should apply to `MessageIntermediateCatchEvent` and `MessageBoundaryEvent` > extend `tMessage` => `tBuffered-CEPMessage`, so that the `messageRef` can be reused > OR extension to `messageEventDefinition`: `ExplicitSubscriptionMessageEventDefinition` > [`subscriptionQuery`, `subscriptionTime`, `bufferPolicy`]

A buffer shared across multiple instances or events is more complex than a simple single-event-buffer (that one does not require buffer policies). As soon as the `requestEvent` call can be executed multiple times for the same `queryId`, we need to specify the following aspects:

Buffer policies: (widely based on [Ref paper Sankalita]) `RetrievalPolicy`, `ConsumptionPolicy`, `LifetimePolicy` + buffer maximum age (= combination of lifetime policies)

5.2 BUFFERED EVENT HANDLING

Why do we need a buffer to allow early event subscription?

What is the desired functionality of the event buffer? What functionality (API) does it expose? > this could be seen as an extension to the API that is exposed by a standard CEP Platform > standard platform API: `registerQuery(queryString, notificationRecipient) : queryId`, `deleteQuery(queryId)` > extended API: `registerQuery(queryString): queryId`, `requestEvent(queryId, notificationRecipient)`, `unsubscribe(queryId)`, `deleteQuery(queryId)`

5.3 EXTENDED PROCESS ENGINE BEHAVIOUR

=> As a link between the BPMN model and the Buffered Event handling

there must be a "subscription-garbage-collection" for any events that cannot be reached anymore in the current process execution!
e.g. two different events behind an xor-gateway. the garbage collection could be executed on every transition

DECLARATION

Put your declaration here.

Potsdam, August 2017

Dennis Wolf

COLOPHON

This document was typeset using the typographical look-and-feel classicthesis developed by André Miede. The style was inspired by Robert Bringhurst's seminal book on typography "*The Elements of Typographic Style*". classicthesis is available for both \LaTeX and \LyX :

<https://bitbucket.org/amiede/classicthesis/>

Happy users of classicthesis usually send a real postcard to the author, a collection of postcards received so far is featured here:

<http://postcards.miede.de/>